

sklearn


October 22, 2019

```
In [1]: %matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: def logit(x):
        return np.log(x/(1-x))
```

```
In [3]: x=np.arange(0.001,0.999, 0.0001)
y=[logit(n) for n in x]
plt.plot(x,y)
plt.xlabel("Probability")
plt.ylabel("Logit - L")
```


```
Out[3]: Text(0, 0.5, 'Logit - L')
```



```
In [4]: def sigmoid(x):
        return (1/(1+np.exp(-x)))
```

```
In [5]: x=np.arange(-10,10,0.0001)
        y=[sigmoid(n) for n in x]
        plt.plot(x,y)
        plt.xlabel("Logit - L")
        plt.ylabel("Probability")
```

```
Out[5]: Text(0, 0.5, 'Probability')
```



```
In [6]: from sklearn.datasets import load_breast_cancer
        cancer=load_breast_cancer()
```

```
In [7]: X=[]
        for target in range(2):
            X.append([[],[]])
            for i in range(len(cancer.data)):
                if cancer.target[i]==target:
                    X[target][0].append(cancer.data[i][0])
                    X[target][1].append(cancer.data[i][1])


        colours =("r","b")
        fig = plt.figure(figsize(10,8))
        ax=fig.add_subplot(111)
```

```

for target in range(2):
    ax.scatter(X[target][0],
               X[target][1],
               c=colours[target])

ax.set_xlabel("mean radius")
ax.set_ylabel("mean texture")
plt.show()

```



```

In [8]: import matplotlib.pyplot as plt
        from mpl_toolkits.mplot3d import Axes3D
        from sklearn.datasets import load_breast_cancer

In [9]: cancer = load_breast_cancer()
        X=[]
        for target in range(2):
            X.append([[], [], []])
            for i in range(len(cancer.data)):
                if cancer.target[i]==target:
                    X[target][0].append(cancer.data[i][0])

```


```

X[target][1].append(cancer.data[i][1])
X[target][2].append(cancer.data[i][2])

colours = ("r","b")
fig=plt.figure(figsize=(18,15))
ax=fig.add_subplot(111, projection='3d')
for target in range(2):
    ax.scatter(X[target][0],
               X[target][1],
               X[target][2],
               c=colours[target])

ax.set_xlabel("mean radius")
ax.set_ylabel("mean texture")
ax.set_zlabel("mean perimeter")
plt.show()

```



```

In [10]: %matplotlib inline
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches


from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
x=cancer.data[:,0] #mean radius
y=cancer.target#0:malignant, 1:benign
colors = {0:'red', 1:'blue'}
plt.scatter(x,y,
            facecolors='none',
            edgecolors=pd.DataFrame(cancer.target)[0].apply(lambda x:colors[x]),cmap=c
plt.xlabel("mean radius")
plt.ylabel("Result")

red = mpatches.Patch(color='red', label='malignant')
blue=mpatches.Patch(color='blue', label='benign')

plt.legend(handles=[red, blue], loc=1)

```

Out[10]: <matplotlib.legend.Legend at 0x21fdd455ac8>



```
In [11]: from sklearn import linear_model
import numpy as np

log_regress = linear_model.LogisticRegression()

In [12]: # train the model
log_regress.fit(X=np.array(x).reshape(len(x),1),y=y)

# print trained model intercept
print(log_regress.intercept_)
print(log_regress.coef_)
```

```
[8.19393897]
[[-0.54291739]]
```

d:\dev\python\python36\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: L

FutureWarning)

```
In [13]: # Plotting the Sigmoid Curve
def sigmoid(x):
    return (1/(1+np.exp(-(log_regress.intercept_[0]+(log_regress.coef_[0][0]*x)))))

x1=np.arange(0,30,0.01)
```

```


y1=[sigmoid(n) for n in x1]

plt.scatter(x,y,facecolors='none',
            edgecolors=pd.DataFrame(cancer.target)[0].apply(lambda x: colors[x]),
            cmap=colors)

plt.plot(x1,y1)
plt.xlabel("mean radius")
plt.ylabel('Probability')

```

Out[13]: Text(0, 0.5, 'Probability')



```
In [14]: print(log_regress.predict_proba([[20]]))
```

```
[[0.93489354 0.06510646]]
```

```
In [15]: print(log_regress.predict([[20]])[0])
```

```
0
```

```
In [16]: print(log_regress.predict_proba([[8]]))
```

```
[[0.02082411 0.97917589]]
```

```
In [17]: print(log_regress.predict([[8]])[0])
```

1

```
In [18]: # Training the model using all features
from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
```

```
In [19]: from sklearn.model_selection import train_test_split
train_set, test_set, train_labels, test_labels = train_test_split(
    cancer.data,
    cancer.target,
    test_size=0.25,
    random_state=1,
    stratify=cancer.target
)
```

```
In [20]: from sklearn import linear_model
x=train_set[:,0:30]
y=train_labels
log_regress=linear_model.LogisticRegression()
log_regress.fit(X=x, y=y)
```

d:\dev\python\python36\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: L

FutureWarning)

```
Out[20]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False)
```

```
In [21]: print(log_regress.intercept_)
```

[0.34532875]

```
In [22]: print(log_regress.coef_)
```

```
[[ 1.80111966e+00  2.55753177e-01 -3.76203243e-02 -5.88987979e-03
 -9.58049351e-02 -3.16746982e-01 -5.06749680e-01 -2.53240777e-01
 -2.26207024e-01 -1.03696078e-02  4.03711661e-03  9.76796186e-01
  2.02748087e-01 -1.22295425e-01 -8.25625028e-03 -1.41118624e-02
 -5.49936132e-02 -3.33054810e-02 -3.05731116e-02  1.13420163e-04
  1.62877492e+00 -4.35039273e-01 -1.50276583e-01 -2.32832527e-02
 -1.94406863e-01 -9.91538995e-01 -1.42903460e+00 -5.40825444e-01
 -6.28853082e-01 -9.04965298e-02]]
```

```

In [23]: import pandas as pd
         # get the predicted probabilities and convert into a dataframe
         preds_prob=pd.DataFrame(log_regress.predict_proba(X=test_set))
         # assign column names to prediction
         preds_prob.columns = ["Malignant", "Benign"]
         # get the predicted class labels
         preds = log_regress.predict(X=test_set)
         preds_class=pd.DataFrame(preds)
         preds_class.columns = ["Prediction"]

In [24]: # actual diagnosis
         original_result = pd.DataFrame(test_labels)
         original_result.columns = ["Original Result"]

In [25]: # merge the three dataframes into one
         result = pd.concat([preds_prob, preds_class, original_result], axis=1)
         print(result.head(10))

```

	Malignant	Benign	Prediction	Original Result
0	0.999812	1.881729e-04	0	0
1	0.998358	1.642333e-03	0	0
2	0.057984	9.420165e-01	1	1
3	1.000000	9.691544e-08	0	0
4	0.207299	7.927008e-01	1	0
5	0.001227	9.987728e-01	1	1
6	0.096810	9.031903e-01	1	1
7	0.007691	9.923086e-01	1	1
8	1.000000	7.828193e-11	0	0
9	0.057154	9.428460e-01	1	1

```

In [26]: # generate table of predictions vs actual
         print("Confusion Matrix")
         print(pd.crosstab(preds, test_labels))

```

```

Confusion Matrix
col_0  0   1
row_0
0      48   3
1       5  87

```

```

In [27]: from sklearn import metrics
         # view the confusion matrix
         print(metrics.confusion_matrix(y_true = test_labels, # True labels
                                         y_pred = preds))      # Predicted labels

```

```

[[48  5]
 [ 3 87]]

```



```
In [28]: # Computing Accuracy, Recall, Precision, and Other Metrics
print("Accuracy")
print(log_regress.score(X=test_set,
                        y=test_labels))
```

```
Accuracy
0.9440559440559441
```

```
In [29]: # view summary of common classification metrics
print("---Metrics---")
print(metrics.classification_report(
    y_true = test_labels,
    y_pred = preds
))
```

```
---Metrics---
```

	precision	recall	f1-score	support
0	0.94	0.91	0.92	53
1	0.95	0.97	0.96	90
micro avg	0.94	0.94	0.94	143
macro avg	0.94	0.94	0.94	143
weighted avg	0.94	0.94	0.94	143

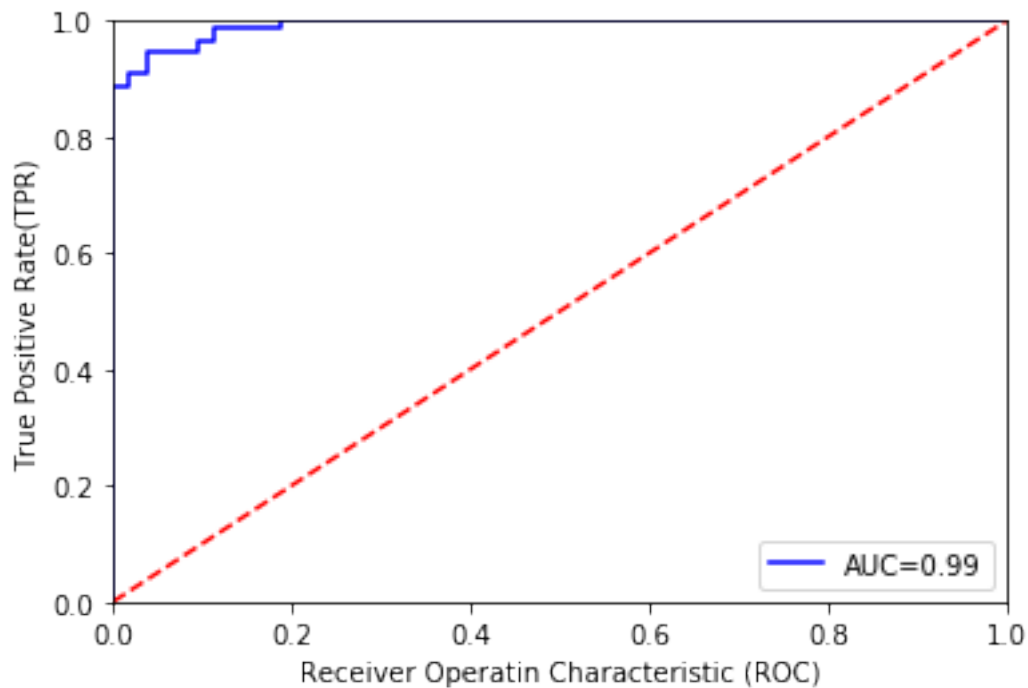
```
In [31]: from sklearn.metrics import roc_curve, auc
#--find the predicted probabilities using the test set
probs = log_regress.predict_proba(test_set)
preds = probs[:,1]

#find the FPR, TPR, and threshold
fpr, tpr, threshold = roc_curve(test_labels, preds)
print(fpr)
print(tpr)
print(threshold)
```

```
[0. 0. 0. 0.01886792 0.01886792 0.03773585
 0.03773585 0.09433962 0.09433962 0.11320755 0.11320755 0.18867925
 0.18867925 1. ]
[0. 0.01111111 0.88888889 0.88888889 0.91111111 0.91111111
 0.94444444 0.94444444 0.96666667 0.96666667 0.98888889 0.98888889
 1. 1. ]
[1.99999109e+00 9.99991091e-01 9.36981948e-01 9.18023512e-01
 9.03190293e-01 8.58497024e-01 8.48205648e-01 5.43404089e-01
 5.25939874e-01 3.71991696e-01 2.71106136e-01 1.21481722e-01
 1.18623350e-01 1.30886736e-21]
```


```
In [32]: #find the area under the curve
roc_auc = auc(fpr, tpr)
```

```
In [33]: import matplotlib.pyplot as plt
plt.plot(fpr, tpr, 'b', label = 'AUC=%0.2f' % roc_auc)
plt.plot([0,1], [0,1], 'r--')
plt.xlim([0,1])
plt.ylim([0,1])
plt.ylabel('True Positive Rate(TPR)')
plt.xlabel('Receiver Operatin Characteristic (ROC)')
plt.legend(loc='lower right')
plt.show()
```



```
In [34]: %matplotlib inline
import pandas as pd
import numpy as np
import seaborn as sns; sns.set(font_scale=1.2)
import matplotlib.pyplot as plt
```

```
In [38]: data=pd.read_csv('svm.csv')
sns.lmplot('x1', 'x2',
           data=data,
           hue='r',
           palette='Set1',
           fit_reg=False,
           scatter_kws={"s": 50});
```



```
In [39]: from sklearn import svm
# Convertin the Columns as Matrices
points = data[['x1', 'x2']].values
result = data['r']

clf = svm.SVC(kernel='linear')
clf.fit(points, result)

print('Vector of weights (w)=', clf.coef_[0])
print('b=', clf.intercept_[0])
print('Indices of support vectors = ', clf.support_)
print('Support vectors = ', clf.support_vectors_)
print('Number of support vectors for each class = ', clf.n_support_)
print('Coefficients of the support vector in the decision function=', np.abs(clf.dual_))
```

Vector of weights (w)= [0.4 0.8]
b= -2.2
Indices of support vectors = [1 2]
Support vectors = [[1. 1.]

[2. 3.]]

Number of support vectors for each class = [1 1]

Coefficients of the support vector in the decision function= [[0.4 0.4]]


```
In [41]: # w is the vector of weights
w = clf.coef_[0]
# find the slope of the hyperplane
slope = -w[0]/w[1]
b = clf.intercept_[0]
# find the coordinates for the hyperplane
xx = np.linspace(0, 4)
yy = slope * xx - (b/w[1])
#plot the margins
s = clf.support_vectors_[0] # first support vector
yy_down = slope * xx + (s[1] - slope * s[0])
s = clf.support_vectors_[-1] # first support vector
yy_up = slope * xx + (s[1] - slope * s[0])

#plot the points
sns.lmplot('x1', 'x2', data =data, hue='r', palette='Set1', fit_reg=False, scatter_kws=
#plot the hyperplane
plt.plot(xx, yy, linewidth=2, color='green');

# plot the 2 margins
plt.plot(xx, yy_down, 'k--')
plt.plot(xx, yy_up, 'k--')
```

```
d:\dev\python\python36\lib\site-packages\ipykernel_launcher.py:11: RuntimeWarning: invalid val
# This is added back by InteractiveShellApp.init_path()
```

```
Out[41]: [<matplotlib.lines.Line2D at 0x21fee52f710>]
```



```
In [44]: print(clf.predict([[3,3]])[0])
          print(clf.predict([[4,0]])[0])
          print(clf.predict([[2,2]])[0])
          print(clf.predict([[1,2]])[0])
```

B
A
B
A



```
In [49]: # Adding a Third Dimension
          %matplotlib inline
          from mpl_toolkits.mplot3d import Axes3D
          import matplotlib.pyplot as plt
          import numpy as np
          from sklearn.datasets import make_circles
          #X is features and c is the class labels
          X, c = make_circles(n_samples = 500, noise = 0.09)
```

```

rgb = np.array(['r', 'g'])
plt.scatter(X[:, 0], X[:, 1], color=rgb[c])
plt.show()

fig = plt.figure(figsize=(18,15))
ax = fig.add_subplot(111, projection='3d')
z = X[:,0]**2 + X[:,1]**2
ax.scatter(X[:, 0], X[:, 1], z, color=rgb[c])
plt.xlabel("x-axis")
plt.ylabel("y-axis")
plt.show()

```



```

In [53]: # combine X (x-axis, y-axis) and z into single ndarray
features = np.concatenate((X,z.reshape(-1,1)), axis=1)

# use SVM for training

```

```

from sklearn import svm

clf = svm.SVC(kernel = 'linear')
clf.fit(features, c)

Out[53]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
          decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
          kernel='linear', max_iter=-1, probability=False, random_state=None,
          shrinking=True, tol=0.001, verbose=False)

In [55]: x3=lambda x,y: (-clf.intercept_[0]-clf.coef_[0][0]*x-clf.coef_[0][1]*y)/clf.coef_[0][2]

In [57]: tmp=np.linspace(-1.5,1.5,100)
          x,y=np.meshgrid(tmp, tmp)

In [59]: ax.plot_surface(x,y,x3(x,y))
          plt.show()

In [64]: from mpl_toolkits.mplot3d import Axes3D
          import matplotlib.pyplot as plt
          import numpy as np
          from sklearn.datasets import make_circles

          #--X is features and c is the class labels
          X, c = make_circles(n_samples=500, noise=0.09)
          z = X[:,0]**2+X[:,1]**2

          rgb = np.array(['r', 'g'])

          fig = plt.figure(figsize=(18,15))
          ax=fig.add_subplot(111, projection='3d')
          ax.scatter(X[:,0], X[:,1], z, color=rgb[c])
          plt.xlabel("x-axis")
          plt.ylabel("y-axis")


          # combine X (x-axis, y-axis) and z into single ndarray
          features = np.concatenate((X,z.reshape(-1,1)), axis=1)

          # use SVM for training
          from sklearn import svm

          clf = svm.SVC(kernel = 'linear')
          clf.fit(features, c)
          x3 = lambda x,y: (-clf.intercept_[0]-clf.coef_[0][0]*x-clf.coef_[0][1]*y)/clf.coef_[0][2]
          tmp = np.linspace(-1.5,1.5,100)
          x,y=np.meshgrid(tmp, tmp)

          ax.plot_surface(x, y, x3(x,y))
          plt.show()

```



```
In [65]: %matplotlib inline
import pandas as pd
import numpy as np
from sklearn import svm, datasets
import matplotlib.pyplot as plt

iris = datasets.load_iris()
print(iris.data[0:5])
print(iris.feature_names)
print(iris.target[0:5])
print(iris.target_names)

[[5.1 3.5 1.4 0.2]
 [4.9 3.  1.4 0.2]
 [4.7 3.2 1.3 0.2]
 [4.6 3.1 1.5 0.2]
 [5.  3.6 1.4 0.2]]
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
[0 0 0 0 0]
['setosa' 'versicolor' 'virginica']


In [84]: X = iris.data[:, :2] # take the first two features
y = iris.target

colors = ['red', 'green', 'blue']
for color, i, target in zip(colors, [0,1,2], iris.target_names):
    plt.scatter(X[y==i, 0], X[y==i, 1], color=color, label=target)

plt.xlabel('Sepal length')
plt.ylabel('Sepal width')
plt.legend(loc='best', shadow=False, scatterpoints=1)

plt.title('Scatter plot of Sepal width against Sepal length')
plt.show()

C = 1 # SVM regularization parameter
clf = svm.SVC(kernel='poly', degree=2, C=C, gamma='auto').fit(X, y)
title = 'SVC with linear kernerl'
```

```
In [85]: # min and max for the first feature
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max()+1
# min and max for the second feature
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max()+1
# step size in the mesh
h = (x_max / x_min)/100
# make predictions for each of the points in xx, yy
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h)
                     )
Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])


# draw the result using a color plot
Z = Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, cmap=plt.cm.Accent, alpha=0.8)

# plot the training points
colors = ['red', 'green', 'blue']
for color, i, target in zip(colors, [0,1,2], iris.target_names):
    plt.scatter(X[y==i, 0], X[y==i, 1], color=color, label=target)

plt.xlabel('Sepal length')
plt.ylabel('Sepal width')
```

```
plt.title(title)
plt.legend(loc='best', shadow=False, scatterpoints=1)
```

Out [85]: <matplotlib.legend.Legend at 0x21fee83e320>



```
In [76]: predictions = clf.predict(X)
         print(np.unique(predictions, return_counts=True))
```

(array([0, 1, 2]), array([50, 53, 47], dtype=int64))

```
In [77]: # Using SVM for Real-Life Problems
```

```
In [86]: #matplotlib inline
import pandas as pd
import numpy as np
from sklearn import svm
import matplotlib.pyplot as plt
import seaborn as sns; sns.set(font_scale=1.2)
```


```
In [87]: data = pd.read_csv('house_sizes_prices_svm.csv')
```

```
In [89]: sns.lmplot('size', 'price',
                   data=data,
```

```

hue='sold',
palette='Set2',
fit_reg=False,
scatter_kws={"s": 50}
);

```



```

In [90]: X = data[['size', 'price']].values
y = np.where(data['sold']=='y', 1, 0) #1 for Y and 0 for N
model = svm.SVC(kernel='linear').fit(X, y)

```

```


In [94]: x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
h = (x_max / x_min) / 20
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h)
                    )

Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
Z=Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, cmap=plt.cm.Blues, alpha=0.3)

```

```
plt.xlabel('Size of house')  
plt.ylabel('Asking price(100s)')  
plt.title("Size of Houses and Their Asking Prices")
```

```
Out[94]: Text(0.5, 1.0, 'Size of Houses and Their Asking Prices')
```



```
In [ ]:
```